

Assignment 2

Course: Social Networks

Name: Mohamed Hassan Ahmed

ID: 2205200

Introduction

This report studies the Facebook Ego Network dataset to detect bots in a social network. We compute graph metrics, train a baseline classifier, and analyze how two attacks—Structural Evasion and Graph Poisoning—affect both the graph structure and detection performance.

1. Import Libraries

Explanation: We import all necessary Python libraries for graph analysis, machine learning, and visualization.

```
[1]: import networkx as nx
import pandas as pd
import numpy as np
from networkx.algorithms.community import label_propagation_communities
from sklearn.ensemble import RandomForestClassifier
from sklearn.model_selection import train_test_split
from sklearn.metrics import accuracy_score, classification_report
import matplotlib.pyplot as plt
```

2. Load Facebook Ego Dataset

Explanation: We load the edges from the dataset and build the network graph.

```
[3]: # Load edges
edges_file = "facebook/0.edges"
edges = pd.read_csv(edges_file, sep=" ", header=None, names=["source", "target"])

# Build graph
G = nx.from_pandas_edgelist(edges, "source", "target")
print(f"Graph loaded with {G.number_of_nodes()} nodes and {G.number_of_edges()} edges")
```

Graph loaded with 333 nodes and 2519 edges

3. Add Node Labels (Simulated Bots)

Explanation: We randomly select 5% of nodes as bots for simulation.

```
[5]: # Simulate some bots (5% of nodes)
num_bots = int(0.05 * G.number_of_nodes())
bot_nodes = np.random.choice(G.nodes(), size=num_bots, replace=False)
labels = {n: 1 if n in bot_nodes else 0 for n in G.nodes()} # 1=bot, 0=normal
```

4. Compute Graph Features

Explanation: We calculate node metrics: degree, clustering coefficient, betweenness centrality, and communities. These are used as features for classification.

```
[7]: # Compute structural features
degree = dict(G.degree())
clustering = nx.clustering(G)
betweenness = nx.betweenness_centrality(G)
communities = list(label_propagation_communities(G))
community_map = {n: i for i, c in enumerate(communities) for n in c}

# Create feature DataFrame
features = pd.DataFrame({
    "node": list(G.nodes()),
    "degree": [degree[n] for n in G.nodes()],
    "clustering": [clustering[n] for n in G.nodes()],
    "betweenness": [betweenness[n] for n in G.nodes()],
    "community": [community_map[n] for n in G.nodes()],
})
features["label"] = features["node"].map(labels)
features.head()
```

```
[7]:   node  degree  clustering  betweenness  community  label
  0    236      36     0.376190    0.003643       0      1
  1    186      43     0.473976    0.005513       0      0
  2    122      62     0.365415    0.022754       0      0
  3    285      46     0.407729    0.013690       0      0
  4     24      15     0.885714    0.000107       1      0
```

5. Baseline Bot Detection

Explanation: Train a Random Forest classifier on the original graph to detect bots. Evaluate using accuracy and classification report.

```
[9]: X = features[["degree", "clustering", "betweenness", "community"]]
y = features["label"]

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)

clf = RandomForestClassifier(n_estimators=100, random_state=42)
clf.fit(X_train, y_train)
y_pred = clf.predict(X_test)

print("== Baseline Bot Detection ==")
print("Accuracy:", accuracy_score(y_test, y_pred))
print(classification_report(y_test, y_pred))
```



```
== Baseline Bot Detection ==
Accuracy: 0.96
          precision    recall  f1-score   support

           0       0.96     1.00     0.98      96
           1       0.00     0.00     0.00       4

      accuracy                           0.96      100
     macro avg       0.48     0.50     0.49      100
weighted avg       0.92     0.96     0.94      100
```

6. Structural Evasion Attack

Explanation: Bots modify their edges—removing some old connections and adding new ones to normal nodes—to blend in and evade detection

```
[11]: G_evasion = G.copy()
      for bot in bot_nodes:
          neighbors = list(G_evasion.neighbors(bot))
          # Remove half of existing edges
          remove_count = max(1, len(neighbors)//2)
          to_remove = np.random.choice(neighbors, remove_count, replace=False)
          G_evasion.remove_edges_from([(bot, n) for n in to_remove])

          # Add 2 new edges to random normal nodes
          normal_nodes = [n for n in G.nodes() if n not in bot_nodes]
          new_edges = np.random.choice(normal_nodes, 2, replace=False)
          for n in new_edges:
              G_evasion.add_edge(bot, n)
```

7. Features After Evasion Attack

Explanation: Recalculate graph features after the evasion attack to retrain the classifier.

```
[13]: degree_ev = dict(G_evasion.degree())
       clustering_ev = nx.clustering(G_evasion)
       betweenness_ev = nx.betweenness_centrality(G_evasion)
       communities_ev = list(label_propagation_communities(G_evasion))
       community_map_ev = {n: i for i, c in enumerate(communities_ev) for n in c}

       features_ev = pd.DataFrame({
           "node": list(G_evasion.nodes()),
           "degree": [degree_ev[n] for n in G_evasion.nodes()],
           "clustering": [clustering_ev[n] for n in G_evasion.nodes()],
           "betweenness": [betweenness_ev[n] for n in G_evasion.nodes()],
           "community": [community_map_ev[n] for n in G_evasion.nodes()],
       })
       features_ev["label"] = features["node"].map(labels)
```

8. Train Classifier After Evasion Attack

Explanation: Train the classifier on the evasion-attacked graph and measure performance drop.

```
[15]: X_ev = features_ev[["degree", "clustering", "betweenness", "community"]]
y_ev = features_ev["label"]

X_train_ev, X_test_ev, y_train_ev, y_test_ev = train_test_split(X_ev, y_ev, test_size=0.3, random_state=42)

clf_ev = RandomForestClassifier(n_estimators=100, random_state=42)
clf_ev.fit(X_train_ev, y_train_ev)
y_pred_ev = clf_ev.predict(X_test_ev)

print("== Bot Detection AFTER Structural Evasion ==")
print("Accuracy:", accuracy_score(y_test_ev, y_pred_ev))
print(classification_report(y_test_ev, y_pred_ev))

== Bot Detection AFTER Structural Evasion ==
Accuracy: 0.96
      precision    recall  f1-score   support
          0       0.96     1.00     0.98      96
          1       0.00     0.00     0.00       4

           accuracy                           0.96      100
          macro avg       0.48     0.50     0.49      100
      weighted avg       0.92     0.96     0.94      100
```

9. Graph Poisoning Attack

Explanation: Bots create many new connections among themselves and to normal users, plus random noise edges in the graph, heavily altering network structure.

```
[17]: G_poison = G.copy()
# Bots connect with each other and extra random nodes
for bot in bot_nodes:
    # Connect to 3 random bots
    bot_partners = np.random.choice(bot_nodes, 3, replace=False)
    for b in bot_partners:
        if bot != b:
            G_poison.add_edge(bot, b)
    # Connect to 5 normal nodes
    normal_nodes = [n for n in G.nodes() if n not in bot_nodes]
    normal_partners = np.random.choice(normal_nodes, 5, replace=False)
    for n in normal_partners:
        G_poison.add_edge(bot, n)
# Add 30 extra random edges between normal nodes
normal_nodes = [n for n in G.nodes() if n not in bot_nodes]
for _ in range(30):
    u, v = np.random.choice(normal_nodes, 2, replace=False)
    G_poison.add_edge(u, v)
```

10. Features After Poisoning Attack

Explanation: Recalculate all features after the poisoning attack for classifier retraining.

```
[19]: degree_p = dict(G_poison.degree())
clustering_p = nx.clustering(G_poison)
betweenness_p = nx.betweenness_centrality(G_poison)
communities_p = list(label_propagation_communities(G_poison))
community_map_p = {n: i for i, c in enumerate(communities_p) for n in c}

features_p = pd.DataFrame({
    "node": list(G_poison.nodes()),
    "degree": [degree_p[n] for n in G_poison.nodes()],
    "clustering": [clustering_p[n] for n in G_poison.nodes()],
    "betweenness": [betweenness_p[n] for n in G_poison.nodes()],
    "community": [community_map_p[n] for n in G_poison.nodes()],
})
features_p["label"] = features["node"].map(labels)
```

11. Train Classifier After Poisoning Attack

Explanation: Retrain the classifier on the poisoned graph to see the performance drop caused by heavy structural changes.

```
[21]: X_p = features_p[["degree", "clustering", "betweenness", "community"]]
y_p = features_p["label"]

X_train_p, X_test_p, y_train_p, y_test_p = train_test_split(X_p, y_p, test_size=0.3, random_state=42)

clf_p = RandomForestClassifier(n_estimators=100, random_state=42)
clf_p.fit(X_train_p, y_train_p)
y_pred_p = clf_p.predict(X_test_p)

print("== Bot Detection AFTER Graph Poisoning ==")
print("Accuracy:", accuracy_score(y_test_p, y_pred_p))
print(classification_report(y_test_p, y_pred_p))

== Bot Detection AFTER Graph Poisoning ==
Accuracy: 0.98
      precision    recall  f1-score   support
          0       0.98     1.00     0.99      96
          1       1.00     0.50     0.67       4
      accuracy         0.98     0.75     0.83     100
    macro avg       0.99     0.75     0.83     100
weighted avg       0.98     0.98     0.98     100
```

12. Visualize Graph Changes

Explanation: Compare graph structure visually before attacks, after structural evasion, and after poisoning. Bots are red; normal nodes are blue.

```
[23]: plt.figure(figsize=(14,6))
pos = nx.spring_layout(G, seed=42)

# Before attack
plt.subplot(1,3,1)
nx.draw(G, pos, node_color=["red" if n in bot_nodes else "skyblue" for n in G.nodes()], node_size=40)
plt.title("Original Graph")

# After Structural Evasion
plt.subplot(1,3,2)
nx.draw(G_evasion, pos, node_color=["red" if n in bot_nodes else "skyblue" for n in G_evasion.nodes()], node_size=40)
plt.title("After Evasion Attack")

# After Poisoning
plt.subplot(1,3,3)
nx.draw(G_poison, pos, node_color=["red" if n in bot_nodes else "skyblue" for n in G_poison.nodes()], node_size=40)
plt.title("After Poisoning Attack")

plt.show()
```

Original Graph After Evasion Attack After Poisoning Attack

